

NIKHIL SOOD

Forward Deployed Engineer — Production LLM Systems | Founder, Nikxius

Dubai, UAE · nikhillsood97@gmail.com · +91 95600 57671 · [linkedin.com/in/nikhillsood](https://www.linkedin.com/in/nikhillsood)

Open to relocating to Paris or Munich · Visa sponsorship required (non-EU) · Available to travel

Founder-engineer who has taken multi-tenant LLM systems from zero to production: built Nikxius, a 14-agent AI SaaS platform serving 5 paying customers, backed by an eval-driven quality loop (hallucination score 0.5→0.95, PII detection 0.7→1.0) and fail-closed safety guards. 6+ years across full-stack engineering, applied ML, and end-to-end customer delivery — comfortable owning ambiguous problems from discovery through technical scoping, production rollout, and iteration with users.

EXPERIENCE

Founder & CEO — Nikxius

Apr 2026 – Present

Kludgedge Apex LLP · Dubai, UAE

- Selected as an Entrepreneur-in-Residence at Antler (Apr–Jun 2026) to build Nikxius from zero to one; completed the program and have continued building and scaling the company independently since.
- Own end-to-end delivery of a multi-tenant, agentic AI SaaS platform — discovery, technical scoping, system design, full-stack build, production rollout, and customer iteration — embedding directly with each customer as primary technical point of contact.
- Shipped 14 production agents (research, enrichment, outreach, qualification, memory, approvals) to 5 paying customers under a 15-qualified-discovery-call-per-month delivery SLA.
- Built an eval-driven quality system — six automated evaluators per LLM run (hallucination, PII, prompt-injection, toxicity, bias, correctness) — lifting hallucination scores 0.5→0.95 and PII detection 0.7→1.0; paired with fail-closed guards and human-in-the-loop approval gates before any customer-facing action.
- Codified recurring customer feedback into reusable deployment playbooks and eval criteria — feeding field learnings back into agent behavior and the product roadmap — while cutting platform cost to ~\$3.25/customer/day at a 50% prompt-cache hit rate and securing \$150K Azure + \$100K AWS credits to fund deployment scale.

Founder & CEO — Kludgedge Apex LLP

Jan 2025 – Mar 2026

New Delhi, India

- Founded a cloud & applied-AI consulting firm, owning discovery, technical scoping, solution design, delivery, and handoff across enterprise engagements in ambiguous, fast-moving environments.
- Led an AI-driven audiology deployment for Hamad Medical Corporation, translating clinical and regulatory stakeholder requirements into a scoped, production-ready technical plan.
- Converted proof-of-concept engagements into revenue-generating production services, making scope/speed/quality trade-offs under deadline pressure without losing stakeholder trust.

Data Scientist — Synaps

Feb 2023 – Jan 2025

Remote / Lyon, France

- Partnered with engineering and compliance stakeholders to ship production fraud/KYC ML for high-stakes detection, monitoring, and reporting workflows in a regulated environment.
- Built and operationalized supervised classifiers and anomaly-detection models (Python, scikit-learn, PyTorch, TensorFlow), improving detection accuracy ~12% and cutting manual review workload ~30%.
- Built ETL pipelines and Postgres data models powering stakeholder-facing monitoring dashboards.

Jr. Data Scientist — IBM

Aug 2019 – Jan 2023

Toronto, Canada

- Supported predictive modeling and forecasting for enterprise product and business stakeholders across large-scale datasets.
- Improved forecasting accuracy ~15% and cut reporting turnaround ~40% by translating model output into decision-ready dashboards (Tableau, Power BI).

SKILLS

AI / LLM: LangGraph, LangSmith, RAG, prompt engineering & caching, fine-tuning, agentic systems, eval frameworks, HITL approval design, scikit-learn, PyTorch, TensorFlow, MLflow

Full-Stack: Python, TypeScript/JavaScript, Next.js, NestJS, REST APIs, Prisma, PostgreSQL, Redis, BullMQ

Cloud & Infra: Azure Container Apps, ACR, AWS, multi-tenant SaaS architecture, CI/CD, cost optimization

Data: SQL, ETL pipelines, feature engineering, model monitoring & retraining, Tableau, Power BI

EDUCATION

MSc, Data Science — University of Toronto

Sep 2017 – Apr 2019

Toronto, Canada · Coursework: statistical modeling, machine learning, deep learning, big data systems, model evaluation

LANGUAGES

English — Native (C2)